ADARSH M

Bangalore, India | adarshm0602@gmail.com | +91 8861550064
LinkedIn | GitHub | Portfolio

SUMMARY

Computer Science graduate with three months of forward deployed engineering internship experience at 1 Hat. Built and improved the Raya triage-engine prototype using Python, LLM workflows and durable storage. Project experience in RAG evaluation, NLP and full-stack web development.

EXPERIENCE

1 Hat | Associate Forward Deployed Engineer Intern

Aug 2026 – Present | Internship scheduled to end 31 Oct 2026

- Developed the Raya triage-engine prototype with Python and FastAPI, integrating LLM-based information extraction, symptom pathways and follow-up question workflows.
- Reduced median message-to-next-response API latency from 17.0 s to 8.5 s across 16 synthetic test turns per profile through row-oriented SQLite storage, caching and deterministic handling of simple answers.
- Repaired uncertainty and numeric-answer handling to preserve evidence and avoid repeated questions; added regression coverage and verified uncertainty handling in two live synthetic journeys.
- Collaborated on kiosk interaction-layer integration, API contracts and reproducible test evidence; kept prototype verification separate from clinical release approval.

TECHNICAL SKILLS

Languages: Python, Java | Backend: FastAPI, Flask, Node.js, Express.js, REST APIs

Frontend: React, HTML, CSS | Databases: SQLite, MySQL, MongoDB

AI and ML: RAG, LLM APIs, LangChain, Hugging Face Transformers, scikit-learn, NLP, retrieval evaluation

Tools: Git, GitHub, pytest, Streamlit, Vercel, Render | Fundamentals: DSA, OOP, DBMS, OS

SELECTED PROJECTS

RAG Evaluation Dataset Builder | Python, embeddings | Code

- Built a timestamp-grounded dataset of 20 QA pairs across three difficulty levels from four ML/DL video transcripts; benchmarked retrieval over 123 chunks using Hit@K and MRR.

Multilingual Sentiment Analysis | Python, NLP, Streamlit | Code

- Built a Kannada, Romanized Kannada and English review pipeline with transliteration, translation and DistilBERT; shipped single-review and batch CSV processing in Streamlit.
- Benchmarked classifiers on 1,154 labeled reviews; Linear SVM achieved 83% macro-F1 versus 81% for Naive Bayes using an 80/20 split. Separate 50-review multilingual evaluation achieved 92% accuracy.

Community Members Connect | Python, Groq, multi-agent workflow | Code

- Built a three-agent CLI demo combining TF-IDF similarity and LLM scoring to match members and draft introductions; saved signals, matches and run summaries for reproducibility.

Match My Car | MongoDB, Express.js, React, Node.js | Demo

- Built a car-discovery application with authentication, filtering and responsive UI; deployed the frontend on Vercel and backend on Render.

EDUCATION

Bachelor of Engineering in Computer Science | JSS Science and Technology University, Mysuru

Graduated June 2026 | CGPA: 8.1